

```

In [8]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from sklearn.linear_model import Ridge, Lasso
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Load the dataset
os.chdir("/Users/ayeshasiddiqha/Downloads")
df = pd.read_csv('jamboree_admission.csv')

# Step 1: Exploratory Data Analysis (EDA)

# Checking the basic structure of the data
print(df.info())
print(df.describe())

# Dropping the Serial No. column as it is not relevant
df = df.drop(columns=['Serial No.'])

# Checking for missing values
print(df.isnull().sum())

# Step 1.1: Univariate Analysis – Distribution Plots for Continuous Variables
continuous_vars = ['GRE Score', 'TOEFL Score', 'University Rating',
                   'SOP', 'CGPA', 'Chance of Admit ']

for var in continuous_vars:
    plt.figure(figsize=(8, 6))
    sns.histplot(df[var], kde=True)
    plt.title(f'Distribution of {var}')
    plt.show()

# Step 1.2: Bivariate Analysis – Correlation Matrix
corr_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=1)
plt.title('Correlation Matrix')
plt.show()

# Step 2: Data Preprocessing

# 2.1: Duplicate value check
print(f'Duplicate rows: {df.duplicated().sum()}')

# 2.2: Missing value treatment (Impute with median for continuous variables)
df = df.fillna(df.median())

# 2.3: Outlier treatment – Checking outliers using IQR
Q1 = df[continuous_vars].quantile(0.25)
Q3 = df[continuous_vars].quantile(0.75)
IQR = Q3 - Q1

outlier_condition = ((df[continuous_vars] < (Q1 - 1.5 * IQR)) | (df[conti

```

```

df[outlier_condition] = np.nan
df = df.fillna(df.median())

# 2.4: Feature Engineering – Standardization (Scaling continuous variable
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()

df[['GRE Score', 'TOEFL Score', 'University Rating',
    'SOP', 'CGPA']] = scaler.fit_transform(df[['GRE Score', 'TOEFL Score',
    'University Rating',
    'SOP', 'CGPA']])

# Step 3: Model Building – Linear Regression
# Separate features and target variable
X = df.drop(columns=['Chance of Admit '])
y = df['Chance of Admit ']

# Adding a constant to the model (for intercept)
X = sm.add_constant(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,

# Linear Regression model
model = sm.OLS(y_train, X_train).fit()
print(model.summary())

# Step 4: Testing Assumptions of Linear Regression

# 4.1: Multicollinearity check (VIF)
vif_data = pd.DataFrame()
vif_data['Variable'] = X_train.columns
vif_data['VIF'] = [variance_inflation_factor(X_train.values, i) for i in

# Drop columns with VIF > 5
vif_data = vif_data[vif_data['VIF'] < 5]
print(vif_data)

# 4.2: Mean of residuals
residuals = y_train - model.fittedvalues
print(f'Mean of Residuals: {np.mean(residuals)}')

# 4.3: Linearity of variables (Residual Plot)
plt.scatter(model.fittedvalues, residuals)
plt.axhline(y=0, color='r', linestyle='--')
plt.title('Residuals vs Fitted Values')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.show()

# 4.4: Homoscedasticity (Constant Variance of Errors)
sns.scatterplot(x=model.fittedvalues, y=residuals)
plt.axhline(y=0, color='r', linestyle='--')
plt.title('Homoscedasticity Check')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.show()

# 4.5: Normality of Residuals (Histogram & Q-Q Plot)
sns.histplot(residuals, kde=True)

```

```
plt.title('Residuals Distribution')
plt.show()

import scipy.stats as stats
stats.probplot(residuals, dist="norm", plot=plt)
plt.title('Q-Q Plot for Normality Check')
plt.show()

# Step 5: Model Performance Evaluation
y_pred = model.predict(X_test)

# Metrics
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)
adj_r2 = 1 - (1 - r2) * (len(y_test) - 1) / (len(y_test) - X_test.shape[1])

print(f'Mean Absolute Error (MAE): {mae}')
print(f'Root Mean Squared Error (RMSE): {rmse}')
print(f'R2 Score: {r2}')
print(f'Adjusted R2 Score: {adj_r2}')

# Step 6: Ridge and Lasso Regression
# Ridge Regression
ridge = Ridge(alpha=1.0)
ridge.fit(X_train, y_train)
ridge_pred = ridge.predict(X_test)

# Lasso Regression
lasso = Lasso(alpha=0.1)
lasso.fit(X_train, y_train)
lasso_pred = lasso.predict(X_test)

# Evaluate Ridge and Lasso models
ridge_mae = mean_absolute_error(y_test, ridge_pred)
ridge_rmse = np.sqrt(mean_squared_error(y_test, ridge_pred))
ridge_r2 = r2_score(y_test, ridge_pred)

lasso_mae = mean_absolute_error(y_test, lasso_pred)
lasso_rmse = np.sqrt(mean_squared_error(y_test, lasso_pred))
lasso_r2 = r2_score(y_test, lasso_pred)

print("Ridge Regression Performance:")
print(f'MAE: {ridge_mae}, RMSE: {ridge_rmse}, R2: {ridge_r2}')

print("Lasso Regression Performance:")
print(f'MAE: {lasso_mae}, RMSE: {lasso_rmse}, R2: {lasso_r2}')
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 500 entries, 0 to 499
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Serial No.	500 non-null	int64
1	GRE Score	500 non-null	int64
2	TOEFL Score	500 non-null	int64
3	University Rating	500 non-null	int64
4	SOP	500 non-null	float64
5	LOR	500 non-null	float64
6	CGPA	500 non-null	float64
7	Research	500 non-null	int64
8	Chance of Admit	500 non-null	float64

```
dtypes: float64(4), int64(5)
```

```
memory usage: 35.3 KB
```



```
None
```



	Serial No.	GRE Score	TOEFL Score	University Rating	SOP
\					
count	500.000000	500.000000	500.000000	500.000000	500.000000
mean	250.500000	316.472000	107.192000	3.114000	3.374000
std	144.481833	11.295148	6.081868	1.143512	0.991004
min	1.000000	290.000000	92.000000	1.000000	1.000000
25%	125.750000	308.000000	103.000000	2.000000	2.500000
50%	250.500000	317.000000	107.000000	3.000000	3.500000
75%	375.250000	325.000000	112.000000	4.000000	4.000000
max	500.000000	340.000000	120.000000	5.000000	5.000000

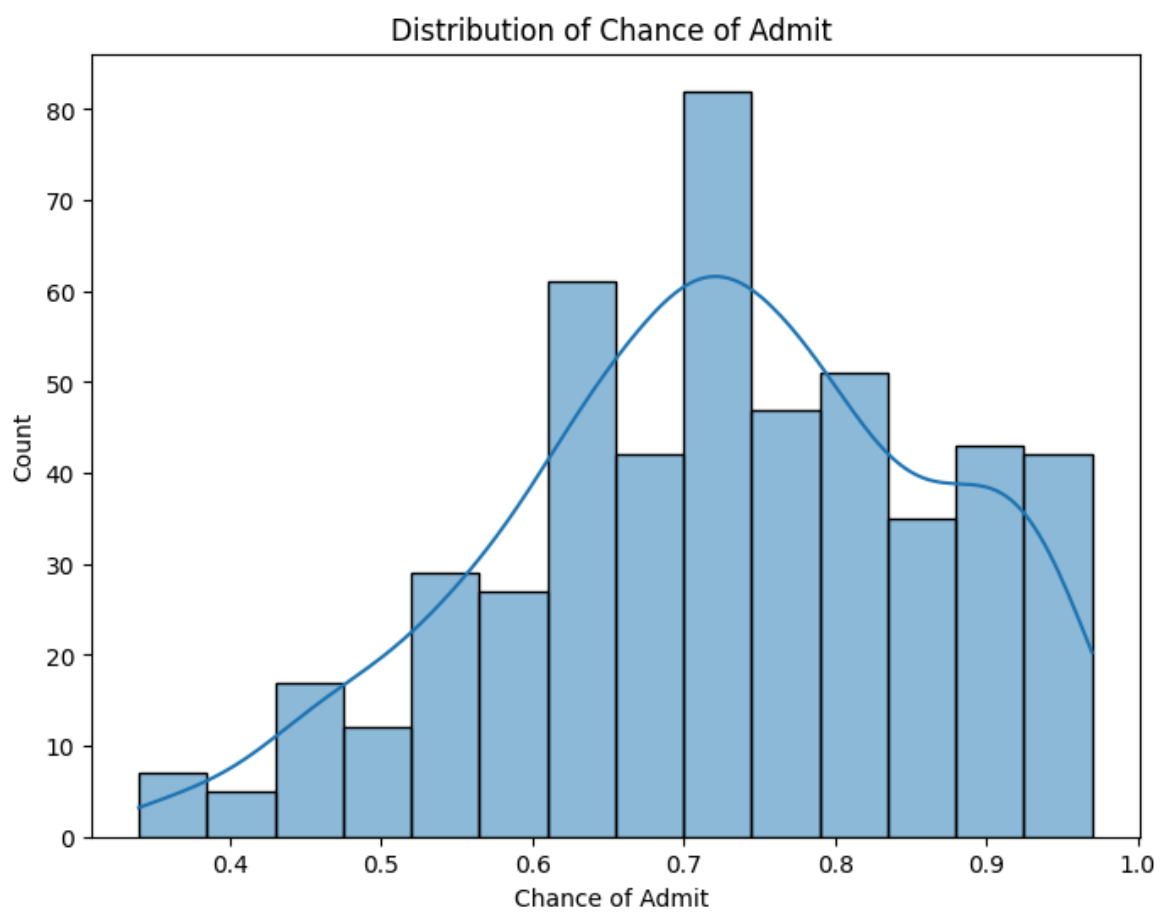

	LOR	CGPA	Research	Chance of Admit
count	500.000000	500.000000	500.000000	500.000000
mean	3.484000	8.576440	0.560000	0.721740
std	0.925450	0.604813	0.496884	0.141140
min	1.000000	6.800000	0.000000	0.340000
25%	3.000000	8.127500	0.000000	0.630000
50%	3.500000	8.560000	1.000000	0.720000
75%	4.000000	9.040000	1.000000	0.820000
max	5.000000	9.920000	1.000000	0.970000


GRE Score	0
TOEFL Score	0
University Rating	0
SOP	0
LOR	0
CGPA	0
Research	0
Chance of Admit	0

```
dtype: int64
```









Duplicate rows: 0

OLS Regression Results

```

=====
=====
Dep. Variable:      Chance of Admit      R-squared:
0.813
Model:              OLS      Adj. R-squared:
0.810
Method:             Least Squares      F-statistic:          2
43.3
Date:               Thu, 12 Dec 2024      Prob (F-statistic):    2.07e
-138
Time:               12:03:39      Log-Likelihood:       56
0.67
No. Observations:   400      AIC:          -1
105.
Df Residuals:       392      BIC:          -1
073.
Df Model:           7
Covariance Type:    nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025
0.975]					

const	0.6598	0.017	39.357	0.000	0.627
0.693					
GRE Score	0.0251	0.007	3.821	0.000	0.012
0.038					
TOEFL Score	0.0155	0.006	2.698	0.007	0.004
0.027					
University Rating	0.0023	0.005	0.484	0.628	-0.007
0.012					
SOP	0.0070	0.005	1.387	0.166	-0.003
0.017					
LOR	0.0150	0.005	3.253	0.001	0.006
0.024					
CGPA	0.0676	0.007	10.353	0.000	0.055
0.080					
Research	0.0231	0.007	3.099	0.002	0.008
0.038					

```

=====
=====
Omnibus:           72.597      Durbin-Watson:
1.974
Prob(Omnibus):     0.000      Jarque-Bera (JB):      18
0.042
Skew:              -0.895      Prob(JB):              8.02
e-40
Kurtosis:          5.756      Cond. No.
22.2
=====
=====

```


Notes:



[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


	Variable	VIF
1	GRE Score	4.489983

2	TOEFL Score	3.664298
3	University Rating	2.572110
4	SOP	2.785764
5	LOR	1.977698
6	CGPA	4.654540
7	Research	1.518065

Mean of Residuals: 3.5735303605122226e-16







Mean Absolute Error (MAE): 0.04181871751792343

Root Mean Squared Error (RMSE): 0.06044901917513528

R2 Score: 0.821316189768417

Adjusted R2 Score: 0.8056077229348713

Ridge Regression Performance:

MAE: 0.04183775276281172, RMSE: 0.06046868388638309, R2: 0.8211999153569034

Lasso Regression Performance:

MAE: 0.10078324542875508, RMSE: 0.1255411111457596, R2: 0.22931195170112628

In []: